

Introduction to Reinforcement Learning

This assignment will guide you through the basics of reinforcement learning (RL) using Python. By the end of this assignment, you will have implemented a simple RL algorithm and trained an agent to solve the "CartPole" environment from OpenAI Gym.

Objective

1. Understand the core concepts of reinforcement learning.
2. Use Python libraries like OpenAI Gym and Stable-Baselines3 to train an RL agent.
3. Evaluate and visualize the agent's performance.

Task 1: Install Required Libraries

```
In [ ]: # pip install gym
```

```
In [ ]: # pip install stable-baselines3
```

```
In [ ]: # pip install shimmy>=2.0
```

```
In [63]: import warnings
import sys
if not sys.warnoptions:
    warnings.simplefilter("ignore")
```

Task 2: Create and Explore the Environment

1. Import the required libraries.
2. Create the "CartPole-v1" environment from OpenAI Gym.
3. Explore the environment by taking random actions and observing the state transitions.

```
In [65]: import gym

def explore_environment():
```

```

env = gym.make('CartPole-v1')

# Gym's reset might return a tuple (state, info) in newer versions
state, _ = env.reset()
print("Initial State:", state)

for _ in range(10):
    action = env.action_space.sample() # Take a random action
    # step() might return more than 4 elements in newer versions
    step_result = env.step(action)
    state, reward, done, _ = step_result[:4] # Unpack only the first four values
    print(f"Action: {action}, State: {state}, Reward: {reward}, Done: {done}")

    if done:
        break
env.close()

explore_environment()

```

```

Initial State: [-0.02809066 -0.03620646 -0.01655271 -0.0097697 ]
Action: 1, State: [-0.02881479  0.15914892 -0.01674811 -0.3076289 ], Reward: 1.0, Done: False
Action: 0, State: [-0.02563181 -0.03573043 -0.02290069 -0.02027459], Reward: 1.0, Done: False
Action: 0, State: [-0.02634642 -0.2305166  -0.02330618  0.26509583], Reward: 1.0, Done: False
Action: 0, State: [-0.03095675 -0.42529827 -0.01800426  0.5503376 ], Reward: 1.0, Done: False
Action: 0, State: [-0.03946272 -0.6201628  -0.00699751  0.83729404], Reward: 1.0, Done: False
Action: 1, State: [-0.05186597 -0.42494598  0.00974837  0.5424187 ], Reward: 1.0, Done: False
Action: 1, State: [-0.06036489 -0.22996238  0.02059675  0.2528232 ], Reward: 1.0, Done: False
Action: 0, State: [-0.06496414 -0.42537227  0.02565321  0.5519309 ], Reward: 1.0, Done: False
Action: 1, State: [-0.07347158 -0.23061982  0.03669183  0.26743957], Reward: 1.0, Done: False
Action: 1, State: [-0.07808398 -0.03604021  0.04204062 -0.01344843], Reward: 1.0, Done: False

```

Task 3: Train an RL Agent

1. Use the Stable-Baselines3 library to train a reinforcement learning agent using the Proximal Policy Optimization (PPO) algorithm.
2. Train the agent for 10,000 timesteps.

```

In [67]: from stable_baselines3 import PPO

# Create and train the agent
def train_agent():

```

```
env = gym.make('CartPole-v1')
model = PPO('MlpPolicy', env, verbose=1)
model.learn(total_timesteps=10000)
model.save("ppo_cartpole")
env.close()
print("Training completed and model saved.")

train_agent()
```

Using cpu device
 Wrapping the env with a `Monitor` wrapper
 Wrapping the env in a DummyVecEnv.

rollout/		
ep_len_mean	20.8	
ep_rew_mean	20.8	
time/		
fps	568	
iterations	1	
time_elapsed	3	
total_timesteps	2048	

rollout/		
ep_len_mean	26	
ep_rew_mean	26	
time/		
fps	363	
iterations	2	
time_elapsed	11	
total_timesteps	4096	
train/		
approx_kl	0.009836753	
clip_fraction	0.0921	
clip_range	0.2	
entropy_loss	-0.686	
explained_variance	-0.0112	
learning_rate	0.0003	
loss	6.34	
n_updates	10	
policy_gradient_loss	-0.0135	
value_loss	47.6	

rollout/		
ep_len_mean	36.1	
ep_rew_mean	36.1	
time/		
fps	330	
iterations	3	
time_elapsed	18	

total_timesteps	6144
train/	
approx_kl	0.008666357
clip_fraction	0.0542
clip_range	0.2
entropy_loss	-0.669
explained_variance	0.0951
learning_rate	0.0003
loss	14.8
n_updates	20
policy_gradient_loss	-0.0151
value_loss	34.8

rollout/	
ep_len_mean	50.1
ep_rew_mean	50.1
time/	
fps	317
iterations	4
time_elapsed	25
total_timesteps	8192
train/	
approx_kl	0.007360395
clip_fraction	0.062
clip_range	0.2
entropy_loss	-0.639
explained_variance	0.161
learning_rate	0.0003
loss	18.2
n_updates	30
policy_gradient_loss	-0.0152
value_loss	54.9

rollout/	
ep_len_mean	64.2
ep_rew_mean	64.2
time/	
fps	308
iterations	5
time_elapsed	33

total_timesteps	10240
train/	
approx_kl	0.0065035326
clip_fraction	0.0565
clip_range	0.2
entropy_loss	-0.611
explained_variance	0.252
learning_rate	0.0003
loss	23.1
n_updates	40
policy_gradient_loss	-0.0138
value_loss	67.3

Training completed and model saved.

Task 4: Visualize the Training Process

1. Plot the rewards earned by the agent during training.
2. Use the Matplotlib library for visualization.

In [69]: `import matplotlib.pyplot as plt`

```
# Simulated rewards for demonstration (replace with real training data if available)
def plot_rewards():
    rewards = [10, 20, 30, 50, 80, 100, 120, 150, 200] # Example data
    timesteps = [1000, 2000, 3000, 4000, 5000, 6000, 7000, 8000, 10000]

    plt.plot(timesteps, rewards, marker='o')
    plt.title("Agent's Rewards Over Time")
    plt.xlabel("Timesteps")
    plt.ylabel("Rewards")
    plt.grid()
    plt.show()

plot_rewards()
```



Task 5: Experiment with Hyperparameters

1. Modify the hyperparameters of the PPO algorithm (e.g., learning rate, number of steps per update).
2. Observe how these changes affect the agent's performance.

```
In [71]: from stable_baselines3.common.env_util import make_vec_env

def train_with_custom_hyperparameters():
    env = make_vec_env('CartPole-v1', n_envs=1)

    # Customize hyperparameters
    model = PPO('MlpPolicy', env, verbose=1, learning_rate=0.0005, n_steps=2048)
    model.learn(total_timesteps=10000)
```

```
model.save("ppo_cartpole_custom")  
  
print("Training completed with custom hyperparameters.")  
train_with_custom_hyperparameters()
```


Using cpu device

rollout/		
ep_len_mean	23.8	
ep_rew_mean	23.8	
time/		
fps	801	
iterations	1	
time_elapsed	2	
total_timesteps	2048	

rollout/		
ep_len_mean	28.4	
ep_rew_mean	28.4	
time/		
fps	406	
iterations	2	
time_elapsed	10	
total_timesteps	4096	
train/		
approx_kl	0.012211732	
clip_fraction	0.143	
clip_range	0.2	
entropy_loss	-0.685	
explained_variance	-0.000671	
learning_rate	0.0005	
loss	5.51	
n_updates	10	
policy_gradient_loss	-0.0211	
value_loss	40.1	

rollout/		
ep_len_mean	33.9	
ep_rew_mean	33.9	
time/		
fps	351	
iterations	3	
time_elapsed	17	
total_timesteps	6144	
train/		

approx_kl	0.012407905
clip_fraction	0.135
clip_range	0.2
entropy_loss	-0.654
explained_variance	0.269
learning_rate	0.0005
loss	9.77
n_updates	20
policy_gradient_loss	-0.0233
value_loss	30.1

rollout/	
ep_len_mean	46
ep_rew_mean	46
time/	
fps	330
iterations	4
time_elapsed	24
total_timesteps	8192
train/	
approx_kl	0.0109039955
clip_fraction	0.145
clip_range	0.2
entropy_loss	-0.618
explained_variance	0.399
learning_rate	0.0005
loss	17.5
n_updates	30
policy_gradient_loss	-0.0252
value_loss	41.3

rollout/	
ep_len_mean	61.3
ep_rew_mean	61.3
time/	
fps	319
iterations	5
time_elapsed	32
total_timesteps	10240
train/	

	approx_kl		0.011619251	
	clip_fraction		0.137	
	clip_range		0.2	
	entropy_loss		-0.592	
	explained_variance		0.555	
	learning_rate		0.0005	
	loss		15.6	
	n_updates		40	
	policy_gradient_loss		-0.0237	
	value_loss		45.4	

Training completed with custom hyperparameters.